

Yullie Yang

yullieyang@gmail.com | github.com/yullieyang | linkedin.com/in/yullie-yang | 240-620-4718

SUMMARY Data scientist and quantitative analyst focused on model validation, experimentation, and applied AI/LLM evaluation in high-stakes settings.

EXPERIENCE

CoStar Group

Boston, MA

Quantitative Analyst

June 2025 - Present

- Supported quarterly releases for commercial real estate credit risk forecasting models, analyzing loan-level datasets to monitor performance and inform model decisions.
- Built a proof-of-concept LLM QA agent and dashboard to support model-monitoring reviews, flag anomalies, and surface evidence for analyst verification.
- Developed Python-based validation workflows and conducted attribution, sensitivity, and stress testing analyses across market and economic scenarios.
- Led release readiness with developers and market analysts by reviewing Git-based code changes, resolving data and logic issues, and validating model outputs.

Guidehouse

McLean, VA

Data Scientist / Consultant

March 2024 - June 2025

- Built Python and SQL analytical workflows across large federal financial datasets, reducing a recurring reconciliation and reporting process from one week to under one day.
- Evaluated RAG and LLM-assisted workflows using Azure AI Search, LangChain, and LangGraph, documenting retrieval issues, failure modes, and implementation criteria.

Federal Reserve Board of Governors, Division of International Finance

Washington, DC

Graduate Intern, Integrated Technology Solutions

May 2023 - August 2023

- Supported data science and application development within the Federal Reserve's International Finance division, contributing to research tools, macroeconomic data workflows, and early GenAI and NLP experimentation.
- Conducted prompt experimentation and NLP-based document analysis to support economic and regulatory research in a policy environment.

PROJECTS

Agentic AI Evaluation & Monitoring Platform | Python · LLM APIs · Streamlit

- Built an LLM-based QA agent that analyzes synthetic model-monitoring data, flags anomalies, and produces evidence-grounded findings for human review.
- Developed an evaluation dashboard to compare prompt and agent configurations across precision, recall, confidence calibration, failure modes, and reviewer decisions.

Product Experimentation & Metrics Analysis | Python · pandas · scipy · statsmodels

- Designed and evaluated a simulated feed-ranking A/B test using 7-day retention as the North Star metric, with SRM checks, power and MDE analysis, CUPED variance reduction, and engagement guardrails.
- Analyzed segment-level treatment effects and retention-engagement tradeoffs, translating results into a data-driven launch recommendation and post-launch monitoring plan.

EDUCATION

University of Pennsylvania — Master of Applied Science in Computer Science

In Progress

University of Maryland — Master of Science in Business Analytics, GPA 3.7

December 2023

Waseda University — Bachelor of Arts in Global Political Economy, GPA 3.7

September 2022

SKILLS

Data & Analytics: Python, SQL, R, Azure, Git, Streamlit, statistical analysis, A/B testing, experimentation, metric design, forecasting, model validation, stress testing

AI/LLM: LLM evaluation, RAG evaluation, prompt experimentation, agent workflows, human-in-the-loop review, retrieval quality analysis, Azure AI Search, LangChain, LangGraph